

Assignment

About Closira

Closira is an AI-powered customer communication platform built for small and medium businesses (SMBs). It handles inbound customer enquiries across WhatsApp, email, and phone, responding using business-defined SOPs (Standard Operating Procedures) to respond accurately, qualify leads, schedule follow-ups, and escalate to human agents when required.

As an AI Engineering Intern, you will be building the intelligence layer that powers these conversations, designing prompts, building agentic workflows, and ensuring the AI behaves safely and reliably in real customer interactions.

Objective

Build a simple AI-powered customer support workflow that demonstrates your ability to design prompts, structure multi-step AI logic, and ensure the system behaves safely and escalates when it should.

Assignment

What you will build

A Python-based AI workflow using the OpenAI or Anthropic Claude API that handles a simulated customer conversation end-to-end across four stages:

#	Stage	What it does
1	FAQ Answering	Answer inbound customer questions using a provided SOP data file (JSON or plain text). Respond only from the SOP; do not hallucinate facts not present in it.
2	Lead Qualification	Ask the customer 2–3 structured questions to qualify them (e.g., business type, team size, current tools). Collect and store responses.
3	Escalation Detection	Detect when the AI should hand off to a human: low confidence, out-of-scope question, angry sentiment, or explicit escalation request. Flag and log the reason.
4	Conversation Summary	At the end of a session, generate a structured summary: customer intent, key details collected, SOP gaps identified, and recommended next action.

SOP Data

Use the following as your SOP source (or create your own for a fictional SMB). Your AI must answer from this data only:

Sample SOP – Closira Demo Business

Business: Bloom Aesthetics Clinic

Hours: Mon-Sat, 9 am-7 pm

Services: Botox (from £200), Fillers (from £250), Consultations (free)

Booking: Via WhatsApp or website. 24hr cancellation required.

Escalate if: complaint, medical question, pricing negotiation, or > 2 unanswered questions.

You may extend or replace this SOP with your own. Clearly document what data your AI operates on.

Prompt Design Requirements

You must submit a prompt_design.md file alongside your code covering:

- **System prompt:** Your full system prompt for the AI, with reasoning for key design choices
- **Hallucination prevention:** How you instruct the model to stay within SOP boundaries and avoid making up information
- **Confidence-based escalation:** How you detect when the AI is uncertain and should hand off (e.g., explicit flag, threshold, output format)
- **Tone and persona:** How you define the AI's communication style for an SMB context

Expected AI Behaviours

Your workflow must handle these scenarios correctly. Include test transcripts in your submission for each:

1. **In-SOP question:** Customer asks, "What are your Botox prices?"
AI answers accurately from SOP only
2. **Out-of-scope question:** Customer asks something not in the SOP.
AI acknowledges the gap and escalates rather than guessing
3. **Escalation trigger:** Customer expresses frustration or a complaint
AI detects sentiment and hands off with a reason
4. **Lead qualification:** AI asks structured questions and produces a qualification summary
5. **Conversation summary:** At session end, AI produces a clean, structured summary with intent, details, and next action

Deliverables

Submit all of the following before the deadline:

- **GitHub repo (public or shared)** containing all code, clearly structured
- **prompt_design.md**, your full system prompt, design decisions, hallucination prevention approach, and escalation logic
- **test_transcripts/ folder**, one sample conversation per expected behaviour listed above (plain text or markdown)
- **README.md**, setup instructions, how to run the workflow, dependencies, and any trade-offs or known limitations
- **2-5 min video** walkthrough explaining your working prototype (code)

Evaluation Criteria

Criterion	What we look for
AI workflow structure	Logical, clean separation of the four stages. Clear flow from intake to summary.
Prompt quality	System prompt is precise, grounded, and well-reasoned. Tone and persona are appropriate for SMB customers.
Reliability & safety	Hallucination prevention is explicit. The AI stays within SOP boundaries and fails gracefully.
Escalation logic	Escalation triggers are clearly defined, reliably detected, and logged with reasons.
SOP understanding	Responses reflect product context. The candidate has thought about real SMB communication needs.
Clarity of reasoning	prompt_design.md is thorough. Decisions are explained, not just implemented.

Submission deadline: 24 hours after receiving the assignment

You do not need to build a full chat UI or frontend. A CLI or script that runs a conversation is sufficient.